
Not All Representational Convergence Is Semantic: Measuring Alignment in the Jacobian Subspace

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The Platonic Representation Hypothesis [12] asserts that neural networks converge
2 towards a shared representation of reality as they become more competent. Evi-
3 dence for this has thus far come from the full activation space, where alignment
4 between two models rises with their competence. However, this comparison mixes
5 semantic content with superficial regularities like document length and syntax.
6 We instead measure alignment in the Jacobian space (J-space), the privileged sub-
7 space induced by the Jacobian lens (J-lens) [10], and find it is a more semantically
8 selective measure.

9 Re-examining convergence by comparing kernels induced by the J-space, we
10 still find that models converge: the Spearman correlation between competence
11 and alignment is $\rho = 0.56$, $p = 0.016$ among language models and $\rho = 0.97$,
12 $p < 10^{-4}$ between language models and vision encoders. Within-language, though,
13 the J-space converges much more slowly than the full activation space, so part of
14 the apparent link between competence and convergence is carried by superficial
15 structure. Consistent with this, the gap between J-space and full convergence
16 disappears across modalities, where no such structure is shared.

17 This convergence is also legible: Using the J-lens at matched layer depth, in-
18 dependently trained models name 5.0 of 25 shared vocabulary items above a
19 position-shuffled floor against 1.9 for the plain logit lens. J-lens overlap converges
20 with competence similarly to the kernel measure, at $\rho = 0.73$, $p = 0.0021$. Code
21 available at <https://github.com/excelexx/jspace-convergence>

22 1 Introduction

23 The Platonic Representation Hypothesis [12] asserts that “Neural networks, trained with different
24 objectives on different data and modalities, are converging to a shared statistical model of reality
25 in their representation spaces.” Concretely, the claim is that alignment between two models grows
26 with their competence. To our knowledge, alignment has thus far been computed over the full
27 representational spaces¹. We show this method conflates semantic content with confounding surface-
28 level similarities like syntax, length and formatting.

29 Recently, Gurnee et al. [10] introduced a new interpretability technique, the Jacobian lens (J-lens),
30 which recovers the concepts a language model is poised to verbalise at a given layer. The J-lens
31 induces the J-space, a privileged subset of internal representations available for report and reasoning.
32 The subspace is both semantically significant, acting as a “global workspace” for the model, and
33 privileged, capturing a median of only 6–7% of a concept vector’s variance [10].

¹We say “full activation space” rather than the “full representational space”, since every quantity we compare is a mean-pooled residual-stream vector at a given layer.

34 Thus, we hypothesise that the J-space will provide a more semantically selective measure of conver-
35 gence than the full activation space, in the sense that J-space alignment will drop more sharply when
36 semantic content is removed.

37 **Motivating experiment** We compute alignment using “mutual κ -nearest-neighbour overlap” be-
38 tween pairs of models. We do this once on an ordinary dataset of text documents and once on the
39 same documents with content words randomised but sentence structure preserved. We find that over
40 half of the full alignment is retained, compared with a quarter of J-space alignment. After running
41 a random-dictionary control, we find that roughly two thirds of this drop in retention is explained
42 by the sparse coding rather than the J-lens dictionary, but the remaining third is attributable to the
43 J-space.

44 We use the J-space to re-examine the Platonic Representation Hypothesis’s central claim, and find
45 that models do indeed converge under this stricter measure.

46 Our three experiments compare 1) J-space against full alignment for text models (including the
47 content-ablated motivating experiment); 2) text models against vision encoders, which have no
48 J-space of their own; and 3) the J-lens readouts of text-models, where outputs are vocabulary tokens
49 and no machinery is needed. For each case, we test convergence by relating alignment to competence.
50 The third experiment also compares against the logit lens [17] to determine whether agreement is
51 attributable to the Jacobian or merely to models predicting similar next words at similar depths.

52 **Related Work** Representational similarity has been measured in many ways: RSA compares
53 dissimilarity matrices [14], SVCCA compares subspaces through canonical correlation [22], and
54 centred kernel alignment compares item–item kernels, with invariance to orthogonal transformation
55 and isotropic scaling [13]. Huh et al. [12] measure convergence using mutual κ -nearest-neighbour
56 overlap [19] and CKNNA.

57 Whether these alignment scores place weight onto semantic content, however, is less settled. Ding
58 et al. [3] show that similarity measures can disagree with functional measures of model behaviour,
59 and validations of these measures have largely focused on whether layers correspond across models,
60 rather than whether such measures isolate meaning.

61 To our knowledge, previous work varies the *metric*, while always comparing the full activation space.
62 In contrast, we hold the metric fixed and restrict the comparison to a specific *subspace*.

63 Since the J-space is defined by the Jacobian to the output, it discards directions that cannot affect the
64 model’s predictions. Therefore, we still compare representations, but unlike previous work, restrict
65 our attention to representations which have functional consequences.

66 Contributions

- 67 1. We propose measuring alignment between models in the J-space rather than the full activation
68 space, and we validate the J-space as a semantically selective measure of alignment using an
69 intervention. We find that randomising content words removes 74% of J-space alignment
70 but only 46% of full-activation alignment.
- 71 2. We find that part of the existing competence–convergence relationship rests on structure
72 that is shared, but not semantic. Indeed, within language, J-space alignment rises with
73 competence at roughly a third the rate of full activation space, while across modalities the
74 two rates are indistinguishable.
- 75 3. We show that J-space convergence is meaningful and can be read as words using the J-
76 lens. At matched relative layer depth, independently trained models name 5.0 of 25 shared
77 vocabulary words against 1.9 for the plain logit lens. J-lens readouts also converge as
78 competence increases, at least as closely as the kernel measure.

79 2 Formal setup

80 We first define the alignment measure used in Experiments 1 and 2. Our third experiment needs no
81 machinery beyond the J-lens itself, since the readouts are in a shared vocabulary space.

For the first two experiments there is no such shared output space. For instance, two models may have different residual widths or different modalities. We therefore adapt the method of Huh et al. [12], which measures alignment by kernels rather than the representations themselves. Between language models, we compare the kernels induced by the J-spaces of the models. Across modalities, we compare the kernel induced by the J-space of the text model to the kernel built from raw pooled patch features.

Let $A, B, \dots$ denote language models with residual widths $r_A, r_B, \dots$ and layer counts $L_A, L_B, \dots$ that generally differ. We index layers by $\ell \in \{0, \dots, L_A - 1\}$ and write $p = \ell / (L_A - 1) \in [0, 1]$ for the corresponding *relative depth*.

We consider alignment over input items. For text–text comparisons, item i is a document and both models receive it. For text–vision comparisons, item i is an image–caption pair, with the vision encoder receiving the image and the language model receiving the caption. For each model, layer ℓ , and item i , we mean-pool the residual stream over the item’s tokens or patches, giving $h_i^{A,(\ell)} \in \mathbb{R}^{r_A}$.

For each language model, we compute the fitted J-lens [10].

$$J_\ell^A = \mathbb{E}_i \left[\mathbb{E}_{(s' \geq s)} \left[\frac{\partial h_{s',i}^{A,(\text{final})}}{\partial h_{s,i}^{A,(\ell)}} \right] \right] \in \mathbb{R}^{r_A \times r_A},$$

where s and s' index tokens with $s' \geq s$, since a token cannot affect its predecessors.

Let $W_U^A \in \mathbb{R}^{m_A \times r_A}$ be model A ’s unembedding matrix, mapping its final residual stream to m_A vocabulary logits. We preprocess W_U^A by folding in the gain w^A of model A ’s final normalisation layer and centring it over the m_A output coordinates:

$$\widetilde{W}_U^A = \left(I_{m_A} - \frac{1}{m_A} \mathbf{1} \mathbf{1}^\top \right) W_U^A \text{diag}(w^A), \quad \mathbf{1} \in \mathbb{R}^{m_A}$$

Without centring, the coefficients for the decomposition in the next step can collapse to near-zero.

The J-space dictionary for layer ℓ of model A , with rows $d_{\ell,v}^A$ that are directions in residual-stream space, is

$$D_\ell^A = \widetilde{W}_U^A J_\ell^A \in \mathbb{R}^{m_A \times r_A}.$$

2.1 Decomposing activations into the J-space and remainder

We decompose each pooled activation into the J-space and non-J remainder, writing each activation as a sparse non-negative combination of $k = 25$ J-space dictionary directions. For tractability, we use a greedy orthogonal matching pursuit (OMP) to select a subset $\mathcal{S}_i$ of dictionary coordinates with $|\mathcal{S}_i| = k$. We then refit the coefficients $\alpha_{i,v}$, where $v \in \mathcal{S}_i$, with a projected-gradient non-negative least-squares (NNLS) step.

Denote the normalised unit directions as $\bar{d}_{\ell,v}^A$ and $h_{\text{full},i}^{A,(\ell)} \equiv h_i^{A,(\ell)}$ as the full pooled activation. Then, we have

$$h_{J,i}^{A,(\ell)} = \sum_{v \in \mathcal{S}_i} \alpha_{i,v} \bar{d}_{\ell,v}^A, \quad h_{\perp,i}^{A,(\ell)} = h_{\text{full},i}^{A,(\ell)} - h_{J,i}^{A,(\ell)}$$

We thus consider three components, written $h_{\omega,i}^{A,(\ell)}$ for $\omega \in \{\text{full}, J, \perp\}$: the full activation $h_{\text{full},i}^{A,(\ell)}$, the J-space component $h_{J,i}^{A,(\ell)}$, and the non-J remainder $h_{\perp,i}^{A,(\ell)}$. Our aim is to compare alignment computed from each.

2.2 Kernels

Each component is winsorised at the 95th percentile and normalised to unit magnitude, and we denote the preprocessed components by $\hat{h}_{\omega,i}^{A,(\ell)}$. The kernel for model A , layer ℓ , and component ω is the $n \times n$ matrix of cosine similarities between items,

$$K_\ell^{A,\omega}(i, j) = \left\langle \hat{h}_{\omega,i}^{A,(\ell)}, \hat{h}_{\omega,j}^{A,(\ell)} \right\rangle.$$

The dictionaries of text models have one row per vocabulary token, making the induced kernels interpretable [2, 8]. However, vision encoders have no such equivalent: MAE [11] decodes to raw pixels, and DINOv2 [18] decodes to abstract prototypes. We thus only apply the J-lens to the caption side, building image kernels from raw pooled patch features, so on the image side, only $\omega = \text{full}$ is defined. The image kernel $K_{\ell}^{\text{img}, \text{full}}$ is formed by the same cosine-similarity expression from the preprocessed² mean-pooled patch activations.

2.3 Mutual κ -nearest-neighbour alignment

We compare two kernels through mutual κ -nearest-neighbour overlap [19]. Taking each item’s κ nearest neighbours under each kernel (the κ largest off-diagonal entries of its row), we define $\text{m-NN}(K_{\ell}^{A, \omega}, K_{\ell'}^{B, \omega'})$ between layer ℓ of model A and layer ℓ' of model B as the average fraction of overlapping neighbours shared by the two kernels, over all n items.

For a text–text comparison both sides carry the same component, $\omega' = \omega$. For a text-to-vision comparison the caption side varies over $\omega \in \{\text{full}, J, \perp\}$ while the image side is fixed at $\omega' = \text{full}$.

Finally, we define the alignment between A and B for component ω as the mean over layer pairs $(\ell, \ell') \in \mathcal{B}_A \times \mathcal{B}_B$, where $\mathcal{B}_A = \{\ell : 0.35 \leq \ell/(L_A - 1) \leq 0.90\}$ is the relative layer depth band:

$$\mathcal{A}^{\omega}(A, B) = \text{mean}_{\ell, \ell'} \text{m-NN}(K_{\ell}^{A, \omega}, K_{\ell'}^{B, \omega'})$$

3 Experiment setup

3.1 Text–text alignment and motivating experiment

We compare the following 11 text models of varying size and architecture: Pythia-70M [1], GPT-2 [20], Gemma-3-270M [7], Qwen3.5-0.8B [24], Gemma-3-1B, Qwen3-1.7B [25], Qwen3.5-2B, Gemma-2-2B [6], Qwen3-4B, Qwen3.5-4B, Gemma-3-4B [7].

For the first experiment, we use the prefitted J-lenses on Neuronpedia [15], and compute alignment across all $\binom{11}{2} = 55$ text-text pairs on 1,000 documents from pile-10k [16], a subset of the Pile [5], following Section 2. We take the first 1,000 rows, truncate each to at most 1,500 characters, then to at most 300 tokens, and mean-pool over token positions. Any difference in tokenizer efficiency is quite marginal, from 3.79 to 4.19 characters per token. Measured over the corpus, the least efficient tokenizer still ingests 94.8% as many characters per document as the most efficient. Neighbourhoods are taken at $\kappa = 10$.

Following the procedure described in the previous section, we evaluate shared semantic content for the J-space, the non-J remainder, and the full activation space over the original corpus. We then repeat the measurement on a content-ablated variant of the same 1,000 documents. We preserve syntactic structure, replacing each non-stopword alphabetic token (i.e. content words) with another drawn from the same log-frequency bin. For example, “Remember that children each progress at their own rate.” becomes “Reduce that self each rates at their own special.” Comparing these two results isolates how much of each component’s alignment is semantic.

We then conduct a random-dictionary control to assess how much of the retention drop comes from the sparse coding of the J-space: we run a Gaussian dictionary through the same content-ablation test and report its retention rate.

Finally, we relate alignment to competence, measuring³ model competence with the HellaSwag benchmark [26]. Appendix B repeats the analysis with different competence axes.

3.2 Text-vision alignment

For our second experiment, we compare DINOv2, MAE, CLIP [21], and SigLIP [27] against the 11 language models for a total of $4 \times 11 = 44$ text-vision model pairs. All four are ViT-Base

²The image side is winsorised per coordinate at the 0.5th and 99.5th percentiles, then normalised to unit magnitude.

³All benchmark results were self-computed to ensure consistency.

Table 1: Alignment on real versus content-ablated documents over 55 model pairs. Retention is the percentage of alignment which survives ablation, and is substantially lower for J-space than for full. The lower block repeats the ablation with the J-lens dictionary replaced by a Gaussian dictionary. Bracketed intervals are 95% percentile bootstrap confidence intervals over the 55 model pairs (2,000 resamples).

	Real alignment	Content-ablated alignment	Retention
<i>Components, J-lens dictionary</i>			
full	0.5577 [0.539, 0.579]	0.3009 [0.290, 0.312]	54.1% [53.01, 55.14]
non-J	0.5057 [0.484, 0.529]	0.2609 [0.252, 0.271]	52.0% [50.73, 53.30]
J	0.4202 [0.409, 0.432]	0.1099 [0.108, 0.112]	26.3% [25.63, 27.02]
<i>J-space under a substituted dictionary</i>			
Gaussian	0.2103 [0.202, 0.220]	0.0719 [0.071, 0.073]	34.7% [33.66, 35.85]

encoders [4] with different pretraining objectives. Alignment is computed as in Section 2 over 1,024 image-caption pairs from the Wikipedia-based Image Text (WIT) dataset [23].

For tractability, we compute mean alignment over a subset of the layers. On the text side, we consider at most⁴ five evenly spaced layers from the band $\mathcal{B}_A$. On the vision side, we consider the final hidden state along with five evenly spaced blocks from index 2 to final - 2. For a model pair, this gives $5 \times 6 = 30$ layer pairs.

We consider a competence axis only on the text side, using the same HellaSwag benchmark.

3.3 Text-text shared J-lens output

For our third experiment, we apply the prefitted J-lenses per model and per layer over 200 documents in the pile-10k dataset [16]. Each document is first truncated to 1,500 characters, then to at most 300 tokens. We read the lens out as words, keeping the top $t = 25$ entries.

$$\text{lens}_\ell^A(h) = \text{softmax}(W_U^A \text{norm}^A(J_\ell^A h)).$$

Since the 11 models do not share a tokenizer, we restrict the J-lens to the 31,548 single-token strings in all 11 vocabularies. We also index positions by character rather than by token, evaluating at 6,000 character positions that fall at a token-end word boundary in all 11 tokenizers.

We then measure how many of the $t = 25$ words are shared between model pairs as a function of relative layer depth $p = \ell / (L_A - 1)$. We report this overlap minus the readout overlap obtained when positions are randomly re-paired. To control for the possibility that the two models are simply predicting similar next words, we compare against the logit lens [17], which reveals what the model would predict if it stopped at a given layer by reading the residual stream directly: $\text{softmax}(W_U^A \text{norm}^A(h))$. As in the first experiment, we then relate the resulting overlap to the mean HellaSwag accuracy of the pair to reevaluate the convergence trend.

All experiments were run on a single RTX 3080 (10 GB) GPU, taking a total of around 40 hours.

4 Results

4.1 Text-text alignment and motivating experiment

When we randomise all content words, the full alignment retains a mean of 54.1% of its original value, while J-space alignment retains only 26.3%⁵. J-space alignment therefore depends more heavily on content than the full activation space.

To test how much of this drop comes from the sparse coding itself, we repeat the ablation experiment with a randomised Gaussian dictionary instead of the J-space dictionary. Under this control, retention is 34.7%, meaning sparse coding accounts for 69.7% of the drop between full and J-space retention. However, the remaining third is still attributable to the J-space itself.

⁴The layer band for Pythia-70M contains only 2 layers. All other text models use five layers.

⁵Retention is the mean of per-pair ratios rather than the ratio of means.




Figure 1: Alignment graphed against HellaSwag accuracy for full activation, J-space, and non-J remainder, with a line-of-best-fit. **(a)** Text–text alignment against competence; p -values are model-label permutation values. The J-space converges much more slowly than the full activation space. **(b)** Alignment averaged over the four vision encoders; error bars show standard error of this mean across the four encoders. Alignment rises with competence for every component in both settings.

191 Mean alignment for the J-space (0.420), the remainder (0.506), and the full activation (0.558) is
 192 between $42.0\times$ and $55.7\times$ chance levels.

193 Comparing alignment between two models and the mean of their competences, we find a positive
 194 correlation across the J-space (Spearman correlation $\rho = 0.56$, $p = 0.016$), full activation ($\rho = 0.77$,
 195 $p = 0.0008$), and remainder ($\rho = 0.87$, $p < 0.0001$) over the 55 pairs (Figure 1(a)). Each model
 196 appears in ten pairs, so the pairs are not independent; p -values thus come from permuting competence
 197 across models while holding alignment values fixed.

198 The rate of convergence in the J-space is about a third of that in the full activation space. Alignment
 199 rises with pair competence at $+0.76 \pm 0.09$ per unit HellaSwag accuracy in the full activation space
 200 against $+0.23 \pm 0.07$ in the J-space.

201 4.2 Comparing text models and vision encoders

202 Alignment between text models and vision encoders for each component is between $6.8\times$ and
 203 $7.4\times$ chance. Alignment correlates strongly and positively with competence across the J-space
 204 ($\rho = 0.97$, $p < 10^{-4}$), the full activation space ($\rho = 0.93$, $p = 0.0001$), and the remainder ($\rho = 0.92$,
 205 $p = 0.0002$). We find that the J-space and full activation convergences are nearly identical.

206 Unlike the text–text comparison, we do not correlate competence with alignment across all 44 text–
 207 vision pairs—since the four encoders differ from one another far more than the text models do, this
 208 would obscure the actual trend. Indeed, mean J-space alignment ranges from 0.053 for MAE to 0.087
 209 for SigLIP, a spread far exceeding the standard deviation of cross-modal alignment across the 11 text
 210 models (0.0134).

211 We instead average alignment within each encoder, then average these four values. We report the four
 212 within-encoder correlations, which range from $+0.95$ to $+0.97$, in Appendix A.

213 Cross-modally, there is no shared superficial structure for a full-space metric to capture, so intuitively,
 214 J-space and full alignment should match, and they do. Considered alongside Section 4.1, where they
 215 diverge, this is further evidence that the within-language gap is explained by superficial structure,
 216 rather than simply a property of the J-space.




Figure 2: **(a)** Mean J-lens top-25 word agreement between language models as a function of each model’s relative layer depth $p = \ell / (L_A - 1)$. Δ is the matched overlap minus the position-shuffled overlap; grids are symmetrised before averaging. **(b)** The shuffle-corrected difference of the J-lens and logit-lens overlap, compared to relative layer depth. $p = 0$ is omitted as the logit lens reads the embedding matrix directly, so dominates by construction. Lines are graphed for same-family pairs and cross-family pairs. The shaded region is ± 1 standard deviation across the model pairs in each family group. The mean delta is positive at every depth shown, and peaks in the middle layers. Same-family pairs have noticeably more agreement than cross-family pairs.

217 4.3 Comparing J-lens shared vocabulary for text models

218 We find that at matched relative layer depth, about 5.1 of 25 words are shared, against 2.0 of 25 for
 219 the plain logit lens. Subtracting off the position-shuffled control, which sits at 0.1 of 25 for the J-lens,
 220 we obtain 5.0 of 25 against 1.9 of 25.

221 We also find that across the 38 cross-family model pairs, a model’s best-matching counterpart layer
 222 sits about 11.0% of the total number of layers away, consistent with Figure 2. Conducting the same
 223 experiment for the logit lens, we find that the J-lens overlap beats the logit-lens overlap in 50/55
 224 pairs, with a median difference of +11.7% overlap, or +2.9 of 25 words. In particular, in the middle
 225 layers $p \in [0.2, 0.5]$, J-lens alignment beats logit-lens alignment by 3–5 $\times$. The logit-lens comparison
 226 rules out the trivial explanation that J-lens alignment is the result of models agreeing on superficial
 227 next-token predictions, providing further evidence that the alignment reflects semantically meaningful
 228 content.

229 We also find that J-lens overlap rises with competence, further validating the results of Section 4.1.
 230 Over the same 55 pairs, $\rho = +0.73$ ($p = 0.0021$), shown in Figure 3. The plain logit lens gives
 231 $\rho = +0.44$, which does not reach significance ($p = 0.085$).

232 This readable form of alignment correlates with competence at least as closely as the kernel form,
 233 providing further evidence to confirm our claim of convergence.

234 4.4 Controls

235 In addition to the random-dictionary control above, we run four further controls:

236 **Random-unembedding-row null** When we replace the J-lens dictionary with random unembed-
 237 ding rows, J-space alignment exceeds the strongest of five seeded draws in all 55 pairs (mean margin
 238 +0.148). J-space alignment thus carries information that arbitrary directions do not. For tractability,
 239 each null dictionary is a random subset of 20,000 unembedding rows rather than the full dictionary.

240 **Lens-fitting corpus** To rule out the possibility that J-space alignment emerges from the shared fitting
 241 corpus, we refit so that no two compared models share one. The effect is immaterial (Appendix C.1).




Figure 3: J-lens top-25 word agreement per model pair against the pair’s mean HellaSwag accuracy, coloured by family composition. Each point is the pair’s median Δ (matched overlap minus position-shuffled overlap) over the eleven relative layer depths. p is a model-label permutation value. Agreement rises with competence.

242 **Shuffling image–caption pairs** Shuffling image–caption pairings in the text-vision alignment
 243 experiments reduces alignment to an average of 0.0096, against a chance level of 0.0098. Since
 244 shuffled pairs share no image–caption relationships, the alignment we measure on the actual pairings
 245 is not simply a result of the metric.

246 **Lens degradation** The J-lens is fitted on ordinary text, so the reduced alignment on the content-
 247 ablated documents in the first experiment may be a result of the J-lens simply working less well on
 248 such documents. We find that degrading the lens cannot reproduce the effect of removing semantic
 249 meaning (Appendix C.2).

250 5 Conclusion and limitations

251 In conclusion, our results suggest that full-space convergence is inflated by surface-level agreement.
 252 Within language (where models share superficial regularities), restricting the comparison to the
 253 J-space slows convergence to a third of the rate. Across modalities, where no such regularities exist,
 254 the two components converge at rates indistinguishable within standard error. Nevertheless, the
 255 original convergence claimed by the Platonic Representation Hypothesis still stands, as the J-spaces
 256 and J-lenses still align more at higher competence.

257 This alignment is meaningful and can be read out as actual words. At the same relative layer depth,
 258 language models’ J-lenses have significant overlap, and this overlap is greater than with the plain logit
 259 lens. Measuring convergence with the J-lens reproduces the result of the first experiment, finding that
 260 more competent models share a greater J-lens overlap.

261 **Limitations** First, we test only 11 text models and 4 vision encoders. The text models span 70M–
 262 4.3B parameters, and all four vision encoders are ViT-Base, so we have no competence axis on the
 263 vision side. Future experiments on larger models would require fitting of new J-lenses, which is
 264 computationally intensive. Furthermore, HellaSwag is near its floor for our smallest models: Pythia-
 265 70M scores 30.8% and GPT-2 scores 38.5% against 25% chance. Thus, competence differences at
 266 the bottom of the range are highly compressed.

267 Second, we show the semantic weight of the J-space through only a single method, the retention
 268 statistic, and the single randomised dictionary control.

269 Third, our cross-modal measurement is asymmetric, as for a given image-caption pair, we compare
 270 the J-space of a text model to the raw visual representation of the vision encoder.

271 Fourth, the J-space is a small part of the full activation space, accounting for an average of 18.3% of
272 the pooled activation’s squared norm across the 11 models. The non-J remainder therefore carries the
273 other 81.7% of the squared norm. Our results concern a minority component of the full activation
274 space, albeit a privileged one, and the remainder is close enough to the full activation that the two
275 behave nearly identically.

276 Fifth, we only test the mutual κ -NN measure of kernel alignment. Other methods, such as CKA,
277 SVCCA, and CKNNA, may produce different results.

278 Finally, the J-lens is fitted per token position, but we apply it to mean-pooled document vectors. This
279 means we are assuming that the dictionary’s validity transfers to those vectors.

280 **Next steps** We would like to strengthen our key claim about the J-space’s semantic weight with
281 more tests on larger and more varied model types.

282 We believe that by leveraging the semantically selective properties of the J-space, we can re-examine
283 many other questions in mechanistic interpretability. We would also welcome a theoretical account
284 of why the J-space should be more semantically selective.

285 Finally, our third experiment compares word overlap rather than activation geometry, which is
286 more functional than representational. Since this paper mostly compares representations, we would
287 welcome future research that uses the J-space to measure functional convergence.

References

- [1] Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shrivanshu Purohit, USVSN Sai Prashanth, Edward Raff, et al. Pythia: A suite for analyzing large language models across training and scaling. *International Conference on Machine Learning*, 2023.
- [2] Guy Dar, Mor Geva, Ankit Gupta, and Jonathan Berant. Analyzing transformers in embedding space. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16124–16170, 2023. doi: 10.18653/v1/2023.acl-long.893.
- [3] Frances Ding, Jean-Stanislas Denain, and Jacob Steinhardt. Grounding representation similarity with statistical testing. In *Advances in Neural Information Processing Systems*, 2021.
- [4] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, et al. An image is worth 16x16 words: Transformers for image recognition at scale, 2021. URL <https://arxiv.org/abs/2010.11929>.
- [5] Leo Gao, Stella Biderman, Sid Black, et al. The Pile: An 800GB dataset of diverse text for language modeling, 2020. URL <https://arxiv.org/abs/2101.00027>.
- [6] Gemma Team. Gemma 2: Improving open language models at a practical size, 2024. URL <https://arxiv.org/abs/2408.00118>.
- [7] Gemma Team. Gemma 3 technical report, 2025. URL <https://arxiv.org/abs/2503.19786>.
- [8] Mor Geva, Avi Caciularu, Kevin Wang, and Yoav Goldberg. Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 30–45, 2022. doi: 10.18653/v1/2022.emnlp-main.3.
- [9] Aaron Gokaslan and Vanya Cohen. OpenWebText corpus. <http://Skylion007.github.io/OpenWebTextCorpus>, 2019.
- [10] Wes Gurnee, Nicholas Sofroniew, Adam Pearce, et al. Verbalizable representations form a global workspace in language models, 2026. URL <https://arxiv.org/abs/2607.15495>.
- [11] Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners, 2021. URL <https://arxiv.org/abs/2111.06377>.
- [12] Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. Position: The Platonic Representation Hypothesis. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 20617–20642. PMLR, 2024.
- [13] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning*, 2019.
- [14] Nikolaus Kriegeskorte, Marieke Mur, and Peter Bandettini. Representational similarity analysis – connecting the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 2:4, 2008.
- [15] Johnny Lin. Neuronpedia: Interactive reference and tooling for analyzing neural networks, 2023. URL <https://www.neuronpedia.org>. Software available from neuronpedia.org.
- [16] Neel Nanda. NeelNanda/pile-10k – datasets at Hugging Face. <https://huggingface.co/datasets/NeelNanda/pile-10k>, 2022.
- [17] nostalgebraist. interpreting GPT: the logit lens. LessWrong, 2020. URL <https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens>.
- [18] Maxime Oquab, Timothée Darcet, Théo Moutakanni, et al. DINOv2: Learning robust visual features without supervision, 2024. URL <https://arxiv.org/abs/2304.07193>.
- [19] Shaul Oron, Tali Dekel, Tianfan Xue, William T. Freeman, and Shai Avidan. Best-buddies similarity—robust template matching using mutual nearest neighbors. *IEEE Transactions on Pattern Analysis & Machine Intelligence*, 40(08):1799–1813, August 2018. ISSN 1939-3539. doi: 10.1109/TPAMI.2017.2737424. URL <https://doi.ieeecomputersociety.org/10.1109/TPAMI.2017.2737424>.
- [20] Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019.

- [21] Alec Radford, Jong Wook Kim, Chris Hallacy, et al. Learning transferable visual models from natural language supervision, 2021. URL <https://arxiv.org/abs/2103.00020>.
- [22] Maithra Raghu, Justin Gilmer, Jason Yosinski, and Jascha Sohl-Dickstein. SVCCA: Singular vector canonical correlation analysis for deep learning dynamics and interpretability. In *Advances in Neural Information Processing Systems*, 2017.
- [23] Krishna Srinivasan, Karthik Raman, Jiecao Chen, Michael Bendersky, and Marc Najork. WIT: Wikipedia-based image text dataset for multimodal multilingual machine learning. *CoRR*, abs/2103.01913, 2021. URL <https://arxiv.org/abs/2103.01913>.
- [24] Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026. URL <https://qwen.ai/blog?id=qwen3.5>.
- [25] An Yang, Anfeng Li, Baosong Yang, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [26] Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a machine really finish your sentence?, 2019. URL <https://arxiv.org/abs/1905.07830>.
- [27] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training, 2023. URL <https://arxiv.org/abs/2303.15343>.

A Aggregation robustness

Table 2 gives the cross-modal correlation separately for each vision encoder. Each encoder has a different alignment level, which is why we consider correlations within encoders rather than pool across all 44 pairs. The relationship holds within every encoder, including MAE, which receives no text supervision at all.

Table 2: Cross-modal alignment against HellaSwag accuracy, computed separately within each vision encoder over the 11 text models. The “mean J” column shows alignment for a given encoder across the 11 text models, showing the differences between the encoders that make pooling the 44 pairs inadvisable. p -values are exact permutation values for $n = 11$.

	mean J	full	non-J	J
DINOv2	0.0742	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.97$ $p < 0.0001$
MAE	0.0535	$\rho = +0.91$ $p = 0.0003$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.95$ $p < 0.0001$
CLIP	0.0764	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.97$ $p < 0.0001$
SigLIP	0.0865	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.97$ $p < 0.0001$

B Competence axes

The main paper measures competence with HellaSwag exclusively. Keeping previous alignment values, we vary the competence scale and derive similar results, shown in Table 3 and Table 4.

Table 3: Spearman rank correlation between within-language alignment and competence, per component. The unit is the model pair ($n = 55$) and a pair’s competence is the mean of its two models’. Because each model appears in ten pairs the pairs are not independent and a standard p -value is invalid; the values shown are model-label permutation p -values, which hold the alignment values fixed and permute competence across models. The alignment values are those of Table 1; only the competence axis varies, here log parameter count and HellaSwag.

	log parameters ($n = 55$)	HellaSwag ($n = 55$)
full	$\rho = +0.79$ $p = 0.0006$	$\rho = +0.77$ $p = 0.0008$
non-J	$\rho = +0.87$ $p = 0.0002$	$\rho = +0.87$ $p < 0.0001$
J	$\rho = +0.60$ $p = 0.0094$	$\rho = +0.56$ $p = 0.0162$

Table 4 repeats the cross-modal comparison of Figure 1(b) against all three competence measures. We also measured competence using language-modelling performance, computed as $1 - \text{bits-per-byte}$

over 4M tokens of OpenWebText [9]. We derive similarly strong results across benchmarks as in the main paper.

We follow Huh et al. [12] in using this measure for the cross-modal comparison only, but for the record, this measure gives a within-language J-space correlation of $\rho = +0.35$ ($p = 0.185$), which is not significant.

Table 4: Spearman rank correlation between cross-modal alignment and competence, per component. Each of the 11 text models contributes one point, its alignment averaged over the four vision encoders and each layer-pair grid reduced by its mean; p -values are exact permutation values for $n = 11$. As in Table 3, only the competence axis varies, here log parameter count, HellaSwag, and 1-bits-per-byte.

	log parameters	HellaSwag	1-bits-per-byte
full	$\rho = +0.84$ $p = 0.0018$	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.96$ $p < 0.0001$
non-J	$\rho = +0.83$ $p = 0.0022$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.98$ $p < 0.0001$
J	$\rho = +0.89$ $p = 0.0005$	$\rho = +0.97$ $p < 0.0001$	$\rho = +0.94$ $p < 0.0001$

C Lens robustness controls

C.1 Lens-fitting control

In the main experiment, J-lenses were all fit on the same corpus. As a result, J-space alignment might be a result of the fact that the lenses were fitted on shared text. Figure 4 recomputes text-text alignment, splitting the fitting corpus into two halves and ensuring the J-lenses of two models in a pair do not use the same half. We also repeat the experiment where two models in a pair share the same half of the fitting corpus. Due to computational constraints, we conduct the control on only 10 pairs (spanning four families) among five models: Pythia-70M, GPT-2, Gemma-3-270M, Qwen3.5-0.8B, and Gemma-3-1B.

We find that corpus-sharing has no measurable impact on alignment. Per-pair differences between the shared-half and crossed-half conditions have a median of $+0.0002$. All three alignment conditions remain far above the random-dictionary null for every pair.




Figure 4: Disjoint lens-fitting control over the 10 pairs among 5 models. Lenses are fitted on disjoint halves of WikiText-103. **Left:** crossed-half against shared-half J-component alignment; points lie on $y = x$ within the ± 0.02 materiality band fixed in advance, so refitting on disjoint halves changes nothing. **Middle:** the three fitting conditions per pair — shared-full corpus, shared-half and crossed-half — all sit far above the per-pair random-dictionary null. **Right:** the per-pair difference Δ between shared-half and crossed-half, median $+0.0002$ against an alignment level of 0.42, every value inside the band.

C.2 Lens-degradation control

The J-lens was fitted only on ordinary text, not content-ablated text, so the reduced alignment on the content-ablated documents may be a result of the J-lens simply working less well on such documents. Indeed, the J-space component accounts for 13.5% of activations on content-ablated text against 18.3% on real text.

384 We construct an instrument-only null by reducing the number of dictionary directions on the original
385 corpus until the J-space component accounts for the same fraction of the activations as it does on
386 content-ablated text. This gives a J-space that is degraded but computed on documents with intact
387 meaning. Alignment between model pairs with the degraded J-space is 0.2924 against 0.1099 for the
388 content-ablated text. Even a five-direction J-space, which accounts for less of the activation than the
389 25-direction J-space on content-ablated text, gives alignment $2.1\times$ higher (0.2326 against 0.1099).
390 Thus, the lens degradation cannot reproduce the effect of removing semantic meaning.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The claims and abstract are properly supported by our analysis, figures, and results.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Key limitations are discussed in the conclusion.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: Our paper does not include any theoretical results. Our formal section setup defines a procedure of comparison between models, but does not derive any new theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Detailed steps of the experiment procedure are included in the Formal Setup and Experiment Setup sections of our paper. As well, code will be released upon acceptance.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We include a detailed procedure for replication and will release code upon acceptance.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Training and test details are all provided in the Experiment Setup sections.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Our experimental section includes confidence intervals, p -values, etc., and figures report error bars. Controls are reported against explicit floors, and are detailed in the controls section.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We list our computation usage in the paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: We have read the code of ethics, and our research conforms to it. Our work involves no human subjects or personal data, and releases no new model.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: Our work has no societal impact. This is a foundational measurement paper on how the internal representations of existing models relate to one another, and we see no path from a more semantically selective alignment to any societal harm.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.

- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No models or data are released, and the only code we use operates on preexisting models and public datasets. There is no risk for misuse.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All models, datasets, and prefitted J-lenses are properly cited, and each is used within the terms of its public release.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: Our only new asset is the experimental design. We release code upon acceptance, and include a README file documenting it.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: No crowdsourcing was used.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: No research with human subjects was conducted.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.

- 702 • Depending on the country in which research is conducted, IRB approval (or equivalent)
703 may be required for any human subjects research. If you obtained IRB approval, you
704 should clearly state this in the paper.
- 705 • We recognize that the procedures for this may vary significantly between institutions
706 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
707 guidelines for their institution.
- 708 • For initial submissions, do not include any information that would break anonymity (if
709 applicable), such as the institution conducting the review.

710 16. Declaration of LLM usage

711 Question: Does the paper describe the usage of LLMs if it is an important, original, or
712 non-standard component of the core methods in this research? Note that if the LLM is used
713 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
714 scientific rigor, or originality of the research, declaration is not required.

715 Answer: [N/A]

716 Justification: This paper studies LLMs and does not use them as a component of the
717 experimental method.

718 Guidelines:

- 719 • The answer [N/A] means that the core method development in this research does not
720 involve LLMs as any important, original, or non-standard components.
- 721 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
722 be described.