
Not All Representational Convergence Is Semantic: Measuring Alignment in the Jacobian Subspace

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The Platonic Representation Hypothesis [11] asserts that neural networks converge towards a shared representation of reality as they become more competent.
2 Current evidence for this hypothesis comes from full activation space alignment,
3 which mixes semantic content with superficial regularities like document length
4 and syntax. We instead measure alignment in the Jacobian space (J-space), the
5 privileged subspace induced by the Jacobian lens (J-lens) [9], and find it is a more
6 semantically selective measure, though a random-dictionary control attributes only
7 part of that selectivity to the J-space itself.

8 Comparing kernels induced by the J-space, we still find that models converge: the
9 text–text Spearman correlation between competence and alignment is $\rho = +0.56$,
10 $p = 0.016$, and the text–vision correlation is $\rho = +0.97$, $p < 10^{-4}$. Within
11 language, though, the J-space converges much more slowly than the full activation
12 space, so part of the apparent link between competence and convergence is carried
13 by superficial structure. Consistent with this, the gap between J-space and full
14 convergence disappears across modalities, where no such structure is shared.

15 This convergence is also legible: using the J-lens at matched layer depth, indepen-
16 dently trained models name 5.1 of 25 shared vocabulary words against 2.0 for the
17 plain logit lens. J-lens overlap rises with competence ($\rho = +0.74$, $p = 0.0020$).
18 Code available at <https://github.com/excelexx/jspace-convergence>.
19

20 1 Introduction

21 The Platonic Representation Hypothesis [11] asserts that “Neural networks, trained with different
22 objectives on different data and modalities, are converging to a shared statistical model of reality
23 in their representation spaces.” Concretely, the claim is that alignment between two models grows
24 with their competence. To our knowledge, alignment has thus far been computed over the full
25 representational spaces.¹ We show this method conflates semantic content with confounding surface-
26 level similarities like syntax, length, and formatting.

27 Recently, Gurnee et al. [9] introduced a new interpretability technique, the Jacobian lens (J-lens),
28 which recovers the concepts a language model is poised to verbalise at a given layer. The J-lens
29 induces the J-space, a privileged subset of internal representations available for report and reasoning.
30 The subspace is both semantically significant, acting as a “global workspace” for the model, and
31 privileged, capturing a median of only 6–7% of a concept vector’s variance [9].

32 Thus, we hypothesise that the J-space will provide a more semantically selective measure of conver-
33 gence than the full activation space, in the sense that J-space alignment will drop more sharply when
34 semantic content is removed.

¹We say “full activation space” rather than “full representational space”, since every quantity we compare is a mean-pooled residual-stream vector at a given layer.

35 **Motivating experiment** We compute alignment using “mutual κ -nearest-neighbour overlap” be-
 36 tween pairs of models. We do this once on an ordinary dataset of text documents and once on the same
 37 documents with content words randomised but sentence structure preserved. We find that over half of
 38 the full alignment is retained, compared with a quarter of J-space alignment. A random-dictionary
 39 control separates how much of this gap is attributable to the J-lens itself as opposed to the sparse
 40 coding effect of decomposition (Section 4.1).

41 We use the J-space to re-examine the Platonic Representation Hypothesis’s central claim, and find
 42 that models do indeed converge under this stricter measure.

43 Our three experiments compare 1) J-space against full alignment for text models (including the
 44 content-ablated motivating experiment); 2) text models against vision encoders, which have no
 45 J-space of their own; and 3) the J-lens readouts of text models, where outputs are vocabulary tokens
 46 and no machinery is needed. For each case, we calculate convergence as the correlation between
 47 alignment and competence. The third experiment also compares J-lens overlap with logit lens
 48 overlap [16] as a control.

49 **Related work** Representational similarity has been measured in many ways: RSA compares dis-
 50 similarity matrices [13], SVCCA compares subspaces through canonical correlation [21], and centred
 51 kernel alignment (CKA) compares item–item kernels, with invariance to orthogonal transformation
 52 and isotropic scaling [12]. Huh et al. [11] measure convergence using mutual κ -nearest-neighbour
 53 overlap [18] and CKNN.

54 Whether these alignment scores place weight onto semantic content, however, is less settled. Further-
 55 more, previous work varies the *metric*, while always comparing the full activation space. In contrast,
 56 we hold the metric fixed and restrict the comparison to a specific *subspace*.

57 Since the J-space is defined by the Jacobian to the output, it discards directions that cannot affect the
 58 model’s predictions. Therefore, we still compare representations but, unlike previous work, restrict
 59 our attention to representations that have functional consequences.

60 Contributions

- 61 1. We propose measuring alignment between models in the J-space rather than the full activation
 62 space, and we validate the J-space as a semantically selective measure of alignment using an
 63 intervention. We find that randomising content words removes 74% of J-space alignment
 64 but only 46% of full-activation alignment, though only about one third of this gap can be
 65 directly attributed to the J-space rather than sparse coding.
- 66 2. We find that part of the existing competence–convergence relationship rests on structure
 67 that is shared, but not semantic. Indeed, within language, J-space alignment rises with
 68 competence at roughly a third the rate of the full activation space, while across modalities
 69 the two rates are indistinguishable.
- 70 3. We show that J-space convergence is meaningful and can be read as words using the J-
 71 lens. At matched relative layer depth, independently trained models name 5.1 of 25 shared
 72 vocabulary words against 2.0 for the plain logit lens. J-lens readouts also converge as
 73 competence increases, at least as closely as the kernel measure.

74 2 Formal setup

75 We first define the alignment measure used in Experiments 1 and 2. Our third experiment needs no
 76 machinery beyond the J-lens itself, since the readouts are in a shared vocabulary space.

77 For the first two experiments there is no such shared output space. For instance, two models may
 78 have different residual widths or different modalities. We therefore adapt the method of Huh et al.
 79 [11], which measures alignment by kernels rather than the representations themselves. Between
 80 language models, we compare the kernels induced by the J-spaces of the models. Across modalities,
 81 we compare the kernel induced by the J-space of the text model to the kernel built from raw pooled
 82 patch features.

Let $A, B, \dots$ denote language models with residual widths $r_A, r_B, \dots$ and layer counts $L_A, L_B, \dots$ that generally differ. We index layers by $\ell \in \{0, \dots, L_A - 1\}$ and write $p = \ell / (L_A - 1) \in [0, 1]$ for the corresponding *relative depth*.

We consider alignment over input items. For text–text comparisons, item i is a document and both models receive it. For text–vision comparisons, item i is an image–caption pair, with each encoder receiving its respective input. For each model, layer ℓ , and item i , we mean-pool the residual stream over the item’s tokens or patches, giving $h_i^{A,(\ell)} \in \mathbb{R}^{r_A}$.

For each language model, we compute the fitted J-lens [9]:

$$J_\ell^A = \mathbb{E}_i \left[\mathbb{E}_{(s' \geq s)} \left[\frac{\partial h_{s',i}^{A,(\text{final})}}{\partial h_{s,i}^{A,(\ell)}} \right] \right] \in \mathbb{R}^{r_A \times r_A},$$

where s and s' index tokens with $s' \geq s$, since a token cannot affect its predecessors.

Let $W_U^A \in \mathbb{R}^{m_A \times r_A}$ be model A ’s unembedding matrix, mapping its final residual stream to m_A vocabulary logits. We preprocess W_U^A by folding in the gain w^A of model A ’s final normalisation layer and centring it over the m_A output coordinates:

$$\widetilde{W}_U^A = \left(I_{m_A} - \frac{1}{m_A} \mathbf{1} \mathbf{1}^\top \right) W_U^A \text{diag}(w^A), \quad \mathbf{1} \in \mathbb{R}^{m_A}.$$

The J-space dictionary for layer ℓ of model A , with rows $d_{\ell,v}^A$ that are directions in residual-stream space, is

$$D_\ell^A = \widetilde{W}_U^A J_\ell^A \in \mathbb{R}^{m_A \times r_A}.$$

2.1 Decomposing activations into the J-space and remainder

We write each pooled activation as a sparse non-negative combination of $k = 25$ J-space dictionary directions, which is how we split the activations into their J-space and non-J remainder. To select the subset $\mathcal{S}_i$ of dictionary coordinates with $|\mathcal{S}_i| = k$, we use a greedy orthogonal matching pursuit (OMP). We then refit the coefficients $\alpha_{i,v}$, where $v \in \mathcal{S}_i$, with a projected-gradient non-negative least-squares (NNLS) step.

Denote by $\bar{d}_{\ell,v}^A$ the normalised unit directions, and write $h_{\text{full},i}^{A,(\ell)} \equiv h_i^{A,(\ell)}$ for the full pooled activation. Then, we have

$$h_{J,i}^{A,(\ell)} = \sum_{v \in \mathcal{S}_i} \alpha_{i,v} \bar{d}_{\ell,v}^A, \quad h_{\perp,i}^{A,(\ell)} = h_{\text{full},i}^{A,(\ell)} - h_{J,i}^{A,(\ell)}.$$

We thus consider three components, written $h_{\omega,i}^{A,(\ell)}$ for $\omega \in \{\text{full}, J, \perp\}$: the full activation $h_{\text{full},i}^{A,(\ell)}$, the J-space component $h_{J,i}^{A,(\ell)}$, and the non-J remainder $h_{\perp,i}^{A,(\ell)}$. Our aim is to compare alignment computed from each.

2.2 Kernels

Each component is winsorised at the 95th percentile and normalised to unit magnitude, and we denote the preprocessed components by $\hat{h}_{\omega,i}^{A,(\ell)}$. The kernel for model A , layer ℓ , and component ω is the $n \times n$ matrix of cosine similarities between items,

$$K_\ell^{A,\omega}(i, j) = \left\langle \hat{h}_{\omega,i}^{A,(\ell)}, \hat{h}_{\omega,j}^{A,(\ell)} \right\rangle.$$

The dictionaries of text models have one row per vocabulary token [2, 7], but a vision analogue would not be as easily interpretable: MAE [10] decodes to raw pixels, and DINOv2 [17] decodes to abstract prototypes. We thus apply the J-lens only to the caption side, building image kernels from raw pooled patch features, so on the image side, only $\omega = \text{full}$ is defined. The image kernel $K_\ell^{\text{img}, \text{full}}$ is formed by the same cosine-similarity expression from the preprocessed² mean-pooled patch activations.

²The image side is winsorised per coordinate at the 0.5th and 99.5th percentiles, then normalised to unit magnitude.

117 2.3 Mutual κ -nearest-neighbour alignment

118 We compare two kernels through mutual κ -nearest-neighbour overlap [18]. Taking each item’s
 119 κ nearest neighbours under each kernel (the κ largest off-diagonal entries of its row), we define
 120 m-NN($K_\ell^{A,\omega}, K_{\ell'}^{B,\omega'}$) between layer ℓ of model A and layer ℓ' of model B as the average fraction
 121 of neighbours shared by the two kernels, over all n items.

122 For a text–text comparison, both sides carry the same component, $\omega' = \omega$. For a text–vision
 123 comparison, the caption side varies over $\omega \in \{\text{full}, J, \perp\}$ while the image side is fixed at $\omega' = \text{full}$.

124 Finally, we define the alignment between A and B for component ω as the mean over layer pairs
 125 $(\ell, \ell') \in \mathcal{B}_A \times \mathcal{B}_B$, where $\mathcal{B}_A = \{\ell : 0.35 \leq \ell/(L_A - 1) \leq 0.90\}$ is the relative layer depth band:

$$\mathcal{A}^\omega(A, B) = \text{mean}_{\ell, \ell'} \text{m-NN}(K_\ell^{A,\omega}, K_{\ell'}^{B,\omega'}).$$

126 3 Experiment setup

127 3.1 Text–text alignment and motivating experiment

128 We compare the following 11 text models of varying size and architecture: Pythia-70M [1], GPT-
 129 2 [19], Gemma-3-270M [6], Qwen3.5-0.8B [23], Gemma-3-1B, Qwen3-1.7B [24], Qwen3.5-2B,
 130 Gemma-2-2B [5], Qwen3-4B, Qwen3.5-4B, and Gemma-3-4B.

131 For the first experiment, we use the prefitted J-lenses on Neuronpedia [14], and compute alignment
 132 across all $\binom{11}{2} = 55$ text–text pairs on 1,000 documents from pile-10k [15], a subset of the Pile [4],
 133 following Section 2. We take the first 1,000 rows, truncate each to at most 1,500 characters, then to
 134 at most 300 tokens, and mean-pool over token positions. Any difference in tokeniser efficiency is
 135 quite marginal, from 3.79 to 4.19 characters per token. Measured over the corpus, the least efficient
 136 tokeniser still ingests 94.8% as many characters per document as the most efficient. Neighbourhoods
 137 are taken at $\kappa = 10$.

138 Following the procedure described in the previous section, we evaluate shared semantic content for
 139 the J-space, the non-J remainder, and the full activation space over the original corpus. We then repeat
 140 the measurement on a content-ablated variant of the same 1,000 documents. We preserve syntactic
 141 structure, replacing each non-stopword alphabetic token (i.e. each content word) with another drawn
 142 from the same log-frequency bin. For example, “Remember that children each progress at their
 143 own rate.” becomes “Reduce that self each rates at their own special.” Comparing these two results
 144 isolates how much of each component’s alignment is semantic.

145 To assess how much of the retention drop comes from the sparse coding of the J-space, we run a
 146 Gaussian dictionary through the same content-ablation test and compute its retention rate.

147 Finally, we relate alignment to competence, measuring competence³ using the HellaSwag bench-
 148 mark [25]. Appendix B repeats the analysis with different competence axes.

149 3.2 Text–vision alignment

150 For our second experiment, we compare DINOv2, MAE, CLIP [20], and SigLIP [26] against the
 151 11 language models for a total of $4 \times 11 = 44$ text–vision model pairs. All four are ViT-Base
 152 encoders [3] with different pretraining objectives. Alignment is computed as in Section 2 over 1,024
 153 image–caption pairs from the Wikipedia-based Image Text (WIT) dataset [22].

154 For tractability, we compute mean alignment over a subset of the layers. On the text side, we consider
 155 at most⁴ five evenly spaced layers from the band $\mathcal{B}_A$. On the vision side, we consider the final hidden
 156 state along with five evenly spaced blocks from index 2 to final $- 2$. For a model pair, this gives at
 157 most $5 \times 6 = 30$ layer pairs.

158 We consider a competence axis only on the text side, using the same HellaSwag benchmark.

³All benchmark results were self-computed to ensure consistency.

⁴The layer band for Pythia-70M contains only two layers. All other text models use five layers.

Table 1: Alignment on real versus content-ablated documents over 55 model pairs. Retention is the content-ablated alignment divided by the original alignment, and is substantially lower for the J-space than for the full space. The lower block repeats the ablation with the J-lens dictionary replaced by a Gaussian dictionary. Bracketed intervals are 95% percentile bootstrap confidence intervals over the 55 model pairs (2,000 resamples).

	Real alignment	Content-ablated alignment	Retention
<i>Components, J-lens dictionary</i>			
full	0.5577 [0.539, 0.579]	0.3009 [0.290, 0.312]	54.1% [53.01, 55.14]
non-J	0.5057 [0.484, 0.529]	0.2609 [0.252, 0.271]	52.0% [50.73, 53.30]
J	0.4202 [0.409, 0.432]	0.1099 [0.108, 0.112]	26.3% [25.63, 27.02]
<i>J-space under a substituted dictionary</i>			
Gaussian	0.2103 [0.202, 0.220]	0.0719 [0.071, 0.073]	34.7% [33.66, 35.85]

3.3 Text–text shared J-lens output

For our third experiment, we apply the prefitted J-lenses per model and per layer over 200 documents in the pile-10k dataset [15]. Each document is first truncated to at most 1,500 characters, then to at most 300 tokens. We read the lens out as words, keeping the top $t = 25$ entries.

$$\text{lens}_\ell^A(h) = \text{softmax}(W_U^A \text{norm}^A(J_\ell^A h)).$$

Since the 11 models do not share a tokeniser, we restrict the J-lens to the 31,548 strings that are a single token in all 11 vocabularies. We also index positions by character rather than by token, evaluating at 6,000 character positions that fall at a token-end word boundary in all 11 tokenisers.

We then measure how many of the $t = 25$ words are shared between model pairs as a function of relative layer depth $p = \ell/(L_A - 1)$. We compare against the logit lens [16], which outputs the next-token prediction if the model stopped at a given layer by reading the residual stream directly: $\text{softmax}(W_U^A \text{norm}^A(h))$. This allows us to rule out the possibility that J-lens agreement is a result of models simply predicting similar next words.

We then relate J-lens overlap to mean HellaSwag accuracy for each model pair to re-evaluate the convergence trend measured in Experiment 1.

All experiments were run on a single RTX 3080 (10 GB) GPU, taking a total of around 40 hours.

4 Results

4.1 Text–text alignment and motivating experiment

When we randomise all content words, the full alignment retains a mean of 54.1% of its original value, while J-space alignment retains only 26.3%⁵ (Table 1). To test how much of this drop comes from the sparse coding itself, we repeat the ablation experiment with a randomised Gaussian dictionary instead of the J-space dictionary. Under this control, retention is 34.7%, meaning sparse coding accounts for 69.7% of the drop between full and J-space retention. However, the remaining third is still attributable to the J-space itself.

J-space alignment therefore depends more heavily on content than the full activation space.

Mean alignment for the J-space (0.420), the non-J remainder (0.506), and the full activation (0.558) is significant, between $42.0\times$ and $55.7\times$ the chance level.

For each model pair, we compare alignment and the mean of their competences. We find a positive correlation across the J-space (Spearman correlation $\rho = +0.56$, $p = 0.016$), full activation ($\rho = +0.77$, $p = 0.0008$), and remainder ($\rho = +0.87$, $p < 0.0001$) (Figure 1(a)). Since pairs are not independent, p -values come from permuting competence across models while holding alignment values fixed.

⁵Retention is the mean of per-pair ratios rather than the ratio of means.




Figure 1: Alignment graphed against HellaSwag accuracy for full activation, J-space, and non-J remainder, with a line of best fit. **(a)** Text–text alignment against competence; p -values are model-label permutation values. The J-space converges much more slowly than the full activation space. **(b)** Alignment averaged over the four vision encoders; error bars show standard error of this mean across the four encoders. Alignment rises with competence for every component in both settings.

190 The rate of convergence in the J-space is about a third of that in the full activation space. Alignment
 191 rises with pair competence at $+0.76 \pm 0.09$ per unit HellaSwag accuracy in the full activation space
 192 against $+0.23 \pm 0.07$ in the J-space.

193 4.2 Text–vision alignment

194 Alignment between text models and vision encoders for each component is significant, between $6.8\times$
 195 and $7.4\times$ the chance level. It correlates strongly and positively with competence across the J-space
 196 ($\rho = +0.97$, $p < 10^{-4}$), the full activation space ($\rho = +0.93$, $p = 0.0001$), and the remainder
 197 ($\rho = +0.92$, $p = 0.0002$) (Figure 1(b)).

198 Unlike the text–text comparison, we do not correlate competence with alignment across all 44 text–
 199 vision pairs—since the four encoders differ from one another far more than the text models do, this
 200 would obscure the actual trend. Indeed, mean J-space alignment ranges from 0.053 for MAE to 0.087
 201 for SigLIP, a spread well in excess of the standard deviation of cross-modal alignment across the 11
 202 text models (0.0134).

203 We instead average alignment within each encoder, then average these four values. We report the four
 204 within-encoder J-space correlations, which range from $+0.95$ to $+0.97$, in Appendix A.

205 Cross-modally, there is no shared superficial structure for a full-space metric to capture, so intuitively,
 206 J-space and full alignment should match, and they do. This is further evidence that the gap between
 207 these components within language is explained by superficial structure rather than by a property of
 208 the J-space.

209 4.3 Text–text shared J-lens output

210 We find that at matched relative layer depth, about 5.1 of 25 words are shared, against 2.0 of 25 for
 211 the plain logit lens. In particular, for middle layers $p \in [0.2, 0.5]$, J-lens alignment beats logit lens
 212 alignment by $4\text{--}5\times$. The logit lens comparison rules out the explanation that J-lens alignment arises
 213 because models might agree on next-token predictions.

214 We also find that across the 38 cross-family model pairs, a model’s best-matching counterpart layer
 215 sits, on average, 11.1% of the layer stack away, consistent with Figure 2.




Figure 2: **(a)** Mean J-lens top-25 word agreement between language models as a function of each model’s relative layer depth $p = \ell / (L_A - 1)$. Grids are symmetrised before averaging. **(b)** The difference between J-lens and logit lens overlap, against relative layer depth. $p = 0$ is omitted as the logit lens reads the embedding matrix directly, and so dominates by construction. Lines are shown for same-family and cross-family pairs. The shaded region is ± 1 standard deviation across the model pairs in each family group. The mean difference is positive at every depth shown, and peaks in the middle layers.




Figure 3: J-lens top-25 median overlap over all matched relative layer depths, against the pair’s mean HellaSwag accuracy. Points are coloured by family composition, and p is a model-label permutation value. Agreement rises with competence.

216 Finally, J-lens overlap rises with competence, further validating the results of Section 4.1. Over the
 217 same 55 pairs, overlap and competence correlate at $\rho = +0.74$ ($p = 0.0020$), shown in Figure 3. The
 218 plain logit lens gives $\rho = +0.41$, which does not reach significance ($p = 0.11$).

219 This readable form of alignment correlates with competence at least as strongly as the kernel form,
 220 corroborating the convergence claim.

221 4.4 Controls

222 In addition to the random-dictionary control above, we run five further controls:

223 **Random-unembedding-row null** When we replace the J-lens dictionary with random unembedding rows, J-space alignment exceeds the strongest of five seeded draws in all 55 pairs (mean margin
224 $+0.167$). This shows that J-space alignment captures information that arbitrary directions do not. The
225 null dictionaries use a random subset of 20,000 unembedding rows rather than the full dictionary due
226 to computational constraints.
227

228 **Lens-fitting corpus** To rule out the possibility that J-space alignment emerges from the shared fitting
229 corpus, we refit so that no two compared models share one. The effect is immaterial (Appendix C.1).

230 **Shuffling image-caption pairs** Shuffling image-caption pairings in the text-vision alignment
231 experiments reduces alignment to an average of 0.0096, against a chance level of 0.0098. Since
232 shuffled pairs share no image-caption relationships, the alignment we measure on the actual pairings
233 is not simply a result of the metric.

234 **Lens degradation** The J-lens is fitted on ordinary text, so the reduced alignment on the content-
235 ablated documents in the first experiment may be a result of the J-lens simply working less well on
236 such documents. We find that degrading the lens cannot reproduce the effect of removing semantic
237 meaning (Appendix C.2).

238 **Position-shuffled floor** In Experiment 3, some of the top-25 lens overlap is coincidental, since both
239 models naturally rank common tokens highly. We compare each model’s top-25 list against the other
240 model’s list at a randomly chosen position in the corpus, which preserves the marginal distributions
241 over words while destroying the correspondence between them. Under this control, overlap is 0.1 of
242 25 words for both the J-lens and logit lens, which is negligible.

243 5 Conclusion and limitations

244 In conclusion, our results suggest that full-space convergence is inflated by surface-level agreement.
245 Within language (where models share superficial regularities), restricting the comparison to the J-
246 space slows convergence to a third of the full-space rate. Across modalities, where no such regularities
247 exist, the J-space and full components converge at rates that are indistinguishable within standard
248 error. Nevertheless, the original convergence claimed by the Platonic Representation Hypothesis still
249 stands, as the J-spaces and J-lenses still align more at higher competence.

250 This alignment is meaningful and can be read out as actual words. At the same relative layer depth,
251 language models’ J-lenses have significant overlap, and this overlap is greater than with the plain logit
252 lens. Measuring convergence with the J-lens reproduces the result of the first experiment, finding that
253 more competent models share a greater J-lens overlap.

254 **Limitations** First, we test only 11 text models and four vision encoders. The text models span
255 70M–4.3B parameters, and all four vision encoders are ViT-Base, so we have no competence axis on
256 the vision side. Future experiments on larger models would require fitting of new J-lenses, which is
257 computationally intensive. Furthermore, HellaSwag is near its floor for our smallest models: Pythia-
258 70M scores 30.8% and GPT-2 scores 38.5% against 25% chance. Thus, competence differences at
259 the bottom of the range are highly compressed.

260 Second, we show the semantic weight of the J-space through only the retention statistic and a single
261 random-dictionary control.

262 Third, our cross-modal measurement is asymmetric, as for a given image-caption pair, we compare
263 the J-space of a text model to the raw visual representation of the vision encoder.

264 Fourth, the J-space is a small part of the full activation space, accounting for an average of 18.3% of
265 the pooled activation’s squared norm across the 11 models. The non-J remainder therefore carries the
266 other 81.7% of the squared norm. Our results concern a minority component of the full activation
267 space, albeit a privileged one, and the remainder is close enough to the full activation that the two
268 behave nearly identically.

269 Fifth, we test only the mutual κ -nearest-neighbour measure of kernel alignment. Other methods, such
270 as CKA, SVCCA, and CKNNA, may produce different results.

271 Finally, the J-lens is fitted per token position, but we apply it to mean-pooled document vectors. This
272 means we are assuming that the dictionary’s validity transfers to those vectors.

273 **Next steps** We would like to strengthen our key claim about the J-space’s semantic weight with
274 more tests on larger and more varied model types.

275 We believe that by leveraging the semantically selective properties of the J-space, we can re-examine
276 many other questions in mechanistic interpretability. We would also welcome a theoretical account
277 of why the J-space should be more semantically selective.

278 Finally, our third experiment compares word overlap rather than activation geometry, a comparison
279 that is more functional than representational. Since this paper mostly compares representations, we
280 would welcome future research that uses the J-space to measure functional convergence.

References

- [1] Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, et al. Pythia: A suite for analyzing large language models across training and scaling. *International Conference on Machine Learning*, 2023.
- [2] Guy Dar, Mor Geva, Ankit Gupta, and Jonathan Berant. Analyzing transformers in embedding space. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16124–16170, 2023. doi: 10.18653/v1/2023.acl-long.893.
- [3] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, et al. An image is worth 16x16 words: Transformers for image recognition at scale, 2021. URL <https://arxiv.org/abs/2010.11929>.
- [4] Leo Gao, Stella Biderman, Sid Black, et al. The Pile: An 800GB dataset of diverse text for language modeling, 2020. URL <https://arxiv.org/abs/2101.00027>.
- [5] Gemma Team. Gemma 2: Improving open language models at a practical size, 2024. URL <https://arxiv.org/abs/2408.00118>.
- [6] Gemma Team. Gemma 3 technical report, 2025. URL <https://arxiv.org/abs/2503.19786>.
- [7] Mor Geva, Avi Caciularu, Kevin Wang, and Yoav Goldberg. Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 30–45, 2022. doi: 10.18653/v1/2022.emnlp-main.3.
- [8] Aaron Gokaslan and Vanya Cohen. OpenWebText corpus. <http://Skylion007.github.io/OpenWebTextCorpus>, 2019.
- [9] Wes Gurnee, Nicholas Sofroniew, Adam Pearce, et al. Verbalizable representations form a global workspace in language models, 2026. URL <https://arxiv.org/abs/2607.15495>.
- [10] Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners, 2021. URL <https://arxiv.org/abs/2111.06377>.
- [11] Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. Position: The Platonic Representation Hypothesis. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 20617–20642. PMLR, 2024.
- [12] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning*, 2019.
- [13] Nikolaus Kriegeskorte, Marieke Mur, and Peter Bandettini. Representational similarity analysis – connecting the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 2:4, 2008.
- [14] Johnny Lin. Neuronpedia: Interactive reference and tooling for analyzing neural networks, 2023. URL <https://www.neuronpedia.org>. Software available from neuronpedia.org.
- [15] Neel Nanda. NeelNanda/pile-10k – datasets at Hugging Face. <https://huggingface.co/datasets/NeelNanda/pile-10k>, 2022.
- [16] nostalgebraist. interpreting GPT: the logit lens. LessWrong, 2020. URL <https://www.lesswrong.com/posts/AcKRB8wDpdAN6v6ru/interpreting-gpt-the-logit-lens>.
- [17] Maxime Oquab, Timothée Darcet, Théo Moutakanni, et al. DINOv2: Learning robust visual features without supervision, 2024. URL <https://arxiv.org/abs/2304.07193>.
- [18] Shaul Oron, Tali Dekel, Tianfan Xue, William T. Freeman, and Shai Avidan. Best-buddies similarity—robust template matching using mutual nearest neighbors. *IEEE Transactions on Pattern Analysis & Machine Intelligence*, 40(08):1799–1813, August 2018. ISSN 1939-3539. doi: 10.1109/TPAMI.2017.2737424. URL <https://doi.ieeecomputersociety.org/10.1109/TPAMI.2017.2737424>.
- [19] Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019.
- [20] Alec Radford, Jong Wook Kim, Chris Hallacy, et al. Learning transferable visual models from natural language supervision, 2021. URL <https://arxiv.org/abs/2103.00020>.

- [21] Maithra Raghu, Justin Gilmer, Jason Yosinski, and Jascha Sohl-Dickstein. SVCCA: Singular vector canonical correlation analysis for deep learning dynamics and interpretability. In *Advances in Neural Information Processing Systems*, 2017.
- [22] Krishna Srinivasan, Karthik Raman, Jiecao Chen, Michael Bendersky, and Marc Najork. WIT: Wikipedia-based image text dataset for multimodal multilingual machine learning. *CoRR*, abs/2103.01913, 2021. URL <https://arxiv.org/abs/2103.01913>.
- [23] Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026. URL <https://qwen.ai/blog?id=qwen3.5>.
- [24] An Yang, Anfeng Li, Baosong Yang, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [25] Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a machine really finish your sentence?, 2019. URL <https://arxiv.org/abs/1905.07830>.
- [26] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training, 2023. URL <https://arxiv.org/abs/2303.15343>.

A Aggregation robustness

Table 2 gives the cross-modal correlation separately for each vision encoder. Each encoder has a different alignment level, which is why we consider correlations within encoders rather than pool across all 44 pairs. The relationship holds within every encoder, including MAE, which receives no text supervision at all.

Table 2: Cross-modal alignment against HellaSwag accuracy, computed separately within each vision encoder over the 11 text models. The “mean J” column shows alignment for a given encoder across the 11 text models, highlighting the differences between encoders that make pooling the 44 pairs inadvisable. p -values are exact permutation values for $n = 11$.

	mean J	full	non-J	J
DINOv2	0.0742	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.97$ $p < 0.0001$
MAE	0.0535	$\rho = +0.91$ $p = 0.0003$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.95$ $p < 0.0001$
CLIP	0.0764	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.97$ $p < 0.0001$
SigLIP	0.0865	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.97$ $p < 0.0001$

B Competence axes

The main paper measures competence with HellaSwag exclusively. Keeping previous alignment values, we vary the competence axis and derive similar results, shown in Table 3 and Table 4.

Table 3: Spearman correlation between within-language alignment and mean competence of the two models, per component. We use both HellaSwag and log parameter count as competence axes. The pairs are not independent, so we use model-label permutation p -values, which hold the alignment values fixed and permute competence across models. The alignment values are the same as in Table 1, and only the competence axis varies.

	log parameters ($n = 55$)	HellaSwag ($n = 55$)
full	$\rho = +0.79$ $p = 0.0006$	$\rho = +0.77$ $p = 0.0008$
non-J	$\rho = +0.87$ $p = 0.0002$	$\rho = +0.87$ $p < 0.0001$
J	$\rho = +0.60$ $p = 0.0094$	$\rho = +0.56$ $p = 0.0162$

Table 4 repeats the cross-modal comparison of Figure 1(b) against all three competence measures. We also measured competence using language-modelling performance, computed as $1 - \text{bits-per-byte}$ over 4M tokens of OpenWebText [8]. We derive similar results to those in the main paper.

We follow Huh et al. [11] in using this measure for the cross-modal comparison only; we note that it gives a within-language J-space correlation of $\rho = +0.35$ ($p = 0.185$), which is not significant.

Table 4: Spearman rank correlation between cross-modal alignment and competence, per component. Each of the 11 text models contributes one point, its alignment averaged over the four vision encoders and each layer-pair grid reduced by its mean; p -values are exact permutation values for $n = 11$. As in Table 3, only the competence axis varies, here log parameter count, HellaSwag, and 1 – bits-per-byte.

	log parameters	HellaSwag	1 – bits-per-byte
full	$\rho = +0.84$ $p = 0.0018$	$\rho = +0.93$ $p = 0.0001$	$\rho = +0.96$ $p < 0.0001$
non-J	$\rho = +0.83$ $p = 0.0022$	$\rho = +0.92$ $p = 0.0002$	$\rho = +0.98$ $p < 0.0001$
J	$\rho = +0.89$ $p = 0.0005$	$\rho = +0.97$ $p < 0.0001$	$\rho = +0.94$ $p < 0.0001$

356 C Lens robustness controls

357 C.1 Lens-fitting control

358 In the main experiment, J-lenses were all fitted on the same corpus (WikiText-103). As a result,
 359 J-space alignment might be a result of the fact that the lenses were fitted on shared text. Figure 4
 360 recomputes text–text alignment, splitting the fitting corpus into two halves and ensuring the J-lenses
 361 of two models in a pair do not use the same half. We also repeat the experiment where two models
 362 in a pair share the same half of the fitting corpus. Due to computational constraints, we conduct
 363 the control on only 10 pairs (spanning three families) among five models: Pythia-70M, GPT-2,
 364 Gemma-3-270M, Qwen3.5-0.8B, and Gemma-3-1B.

365 We find that corpus-sharing has no measurable impact on alignment. Per-pair differences between the
 366 shared-half and crossed-half conditions have a median of +0.0002. All three alignment conditions
 367 remain far above the random-dictionary null for every pair.




Figure 4: Disjoint lens-fitting control over the 10 pairs among five models. Lenses are fitted on disjoint halves of WikiText-103. **Left:** crossed-half against shared-half J-component alignment; points lie on $y = x$ within the ± 0.02 materiality band fixed in advance, so refitting on disjoint halves changes nothing. **Middle:** the three fitting conditions per pair—shared-full corpus, shared-half, and crossed-half—all sit far above the per-pair random-dictionary null. **Right:** the per-pair difference Δ between shared-half and crossed-half, median +0.0002 against an alignment level of 0.42, every value inside the band.

368 C.2 Lens-degradation control

369 The J-lens was fitted only on ordinary text, not content-ablated text, so the reduced alignment on the
 370 content-ablated documents may be a result of the J-lens simply working less well on such documents.
 371 Indeed, the J-space component accounts for 13.5% of the activation’s squared norm on content-ablated
 372 text against 18.3% on real text.

373 We construct an instrument-only null by reducing the number of dictionary directions on the original
 374 corpus until the J-space component accounts for the same fraction of the activations as it does on
 375 content-ablated text. This gives a J-space that is degraded but computed on documents with intact
 376 meaning. Alignment between model pairs with the degraded J-space is 0.2924 against 0.1099 for the
 377 content-ablated text. Even a five-direction J-space, which accounts for less of the activation than the
 378 25-direction J-space on content-ablated text, gives alignment $2.1 \times$ higher (0.2326 against 0.1099).

379 Thus, the lens degradation cannot reproduce the effect of removing semantic meaning.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The claims in the abstract and introduction are supported by our analysis, figures, and results.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Key limitations are discussed in the conclusion.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: Our paper does not include any theoretical results. Our formal setup section defines a procedure for comparing models, but does not derive any new theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Detailed steps of the experiment procedure are included in the Formal Setup and Experiment Setup sections of our paper. In addition, code will be released upon acceptance.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We include a detailed procedure for replication and will release code upon acceptance.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Training and test details are all provided in the Experiment Setup sections.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Our experimental section includes confidence intervals, p -values, etc., and figures report error bars. Controls are reported against explicit floors, and are detailed in the controls section.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We report our compute usage in the paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: We have read the code of ethics, and our research conforms to it. Our work involves no human subjects or personal data, and releases no new model.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: Our work has no direct societal impact. This is a foundational measurement paper on how the internal representations of existing models relate to one another, and we see no path from a more semantically selective alignment to any societal harm.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.

- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No models or data are released, and the only code we use operates on preexisting models and public datasets. There is no risk of misuse.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All models, datasets, and prefitted J-lenses are properly cited, and each is used within the terms of its public release.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: Our only new assets are the experimental code and design, released upon acceptance with a README documenting them.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: No crowdsourcing was used.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: No research with human subjects was conducted.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.

- 691 • Depending on the country in which research is conducted, IRB approval (or equivalent)
692 may be required for any human subjects research. If you obtained IRB approval, you
693 should clearly state this in the paper.
- 694 • We recognize that the procedures for this may vary significantly between institutions
695 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
696 guidelines for their institution.
- 697 • For initial submissions, do not include any information that would break anonymity (if
698 applicable), such as the institution conducting the review.

699 **16. Declaration of LLM usage**

700 Question: Does the paper describe the usage of LLMs if it is an important, original, or
701 non-standard component of the core methods in this research? Note that if the LLM is used
702 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
703 scientific rigor, or originality of the research, declaration is not required.

704 Answer: [N/A]

705 Justification: This paper studies LLMs and does not use them as a component of the
706 experimental method.

707 Guidelines:

- 708 • The answer [N/A] means that the core method development in this research does not
709 involve LLMs as any important, original, or non-standard components.
- 710 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
711 be described.